

Scrypt is Maximally Memory Hard

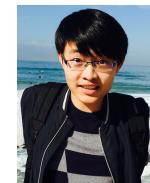
Joël Alwen

IST Austria



Binyi Chen

UCSB



Krzysztof Pietrzak

IST Austria



Leonid Reyzin

Boston University

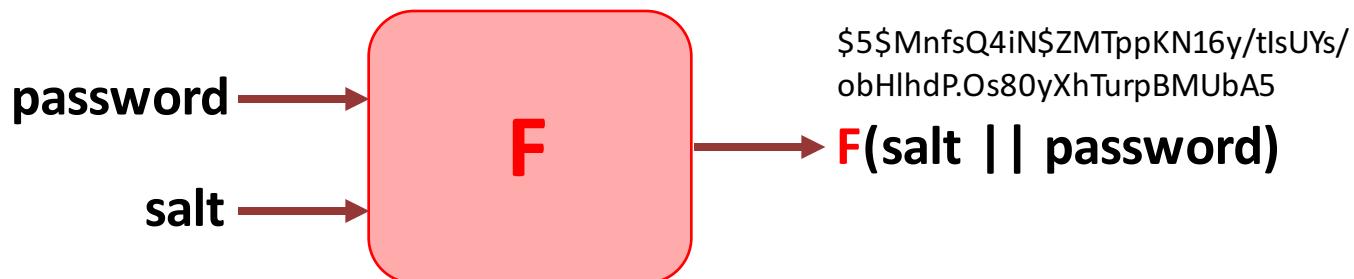


Stefano Tessaro

UCSB



Password hashing: Store a hash of a password + salt



F is moderately hard

Honest users can still login quickly.

≈ 1 evaluation of F

Brute-force attack is infeasible. 

Many evaluations of F

Traditional hardness metric: Time complexity (e.g., PKCS #5)

ASIC-resistance

Honest users

(General-purpose CPU)



cost per **F eval** = C

better parallelization, pipelining for speedup; lower energy costs ...

Adversaries

(ASICs)



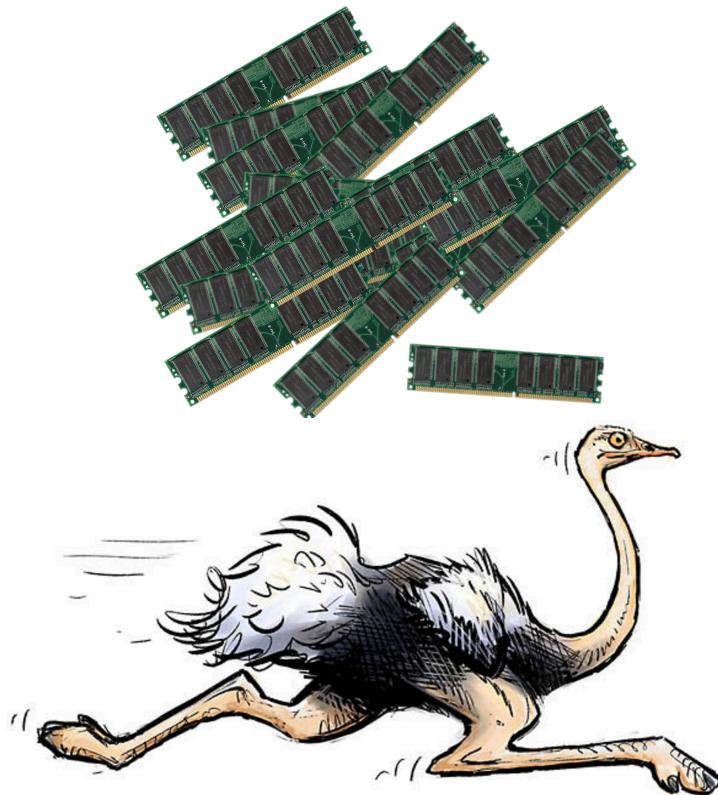
cost per **F eval** = $C' \ll C$



Can we enforce $C' \approx C$?

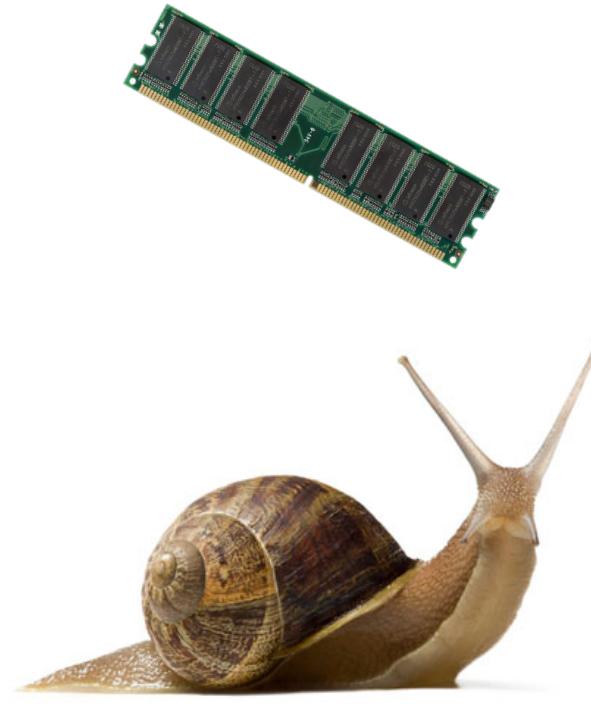
Idea: Memory costs (e.g., on-chip area, access time, \$-cost)
are platform independent

Memory-hard functions (MHFs) [Percival, 2009]



Fast evaluation of MHF F ⇒

large memory

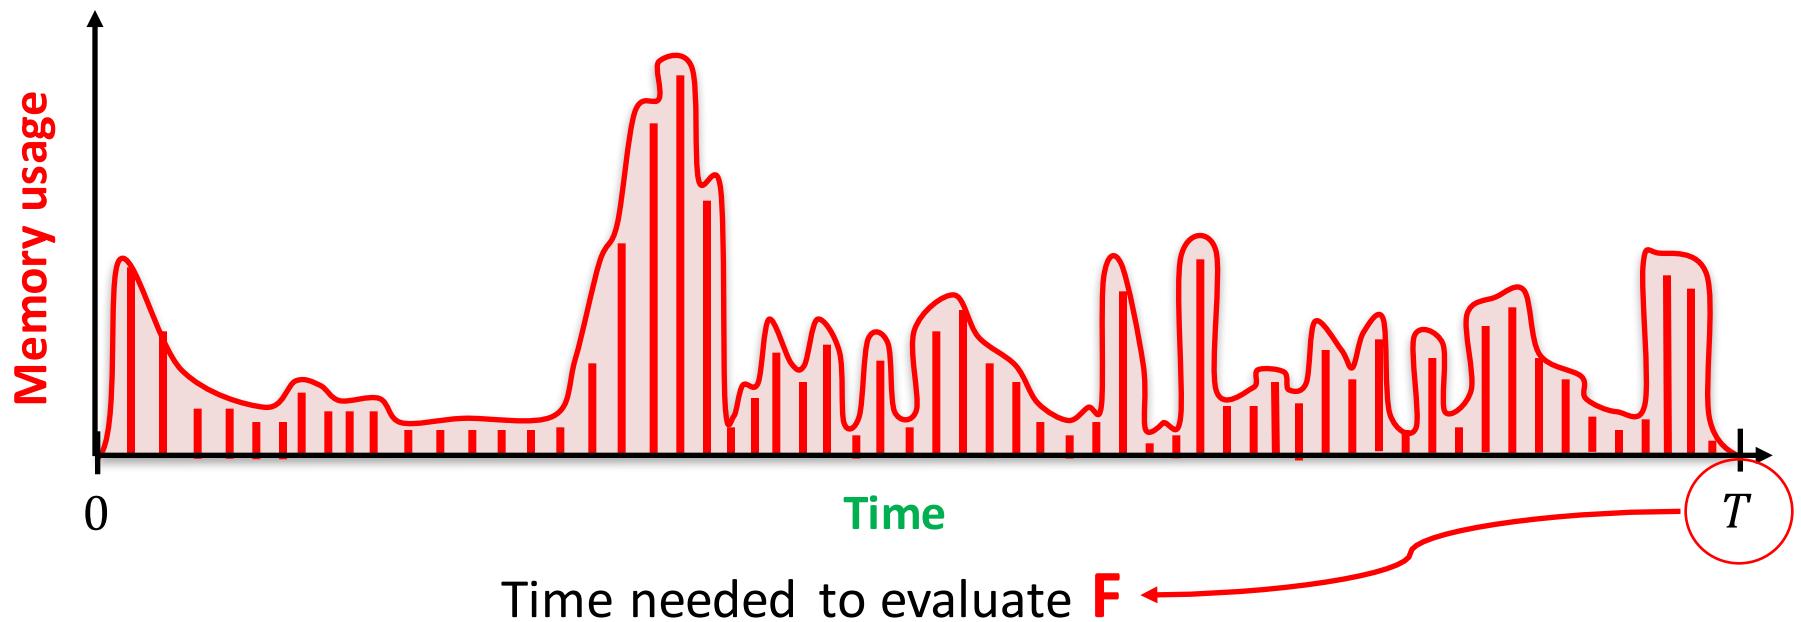


Small memory ⇒

slow evaluation of MHF F

Memory-hardness, more precisely

Goal: Maximize cumulative memory complexity (CMC) for any (possibly **parallelized**) strategy to evaluate F .



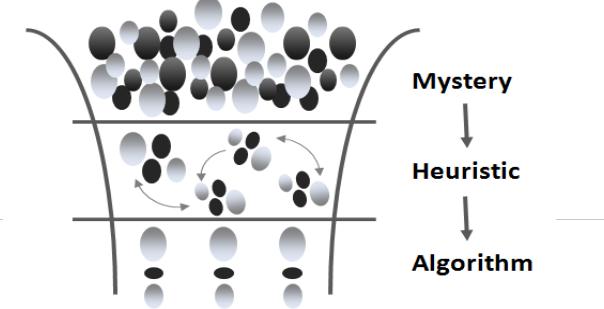
$$\text{CMC} = \sum_{t=0}^T \text{Memory}(t)$$

[Alwen and Serbinenko, STOC '15]

Memory-hardness

PHC call for submissions

The Password Hashing Competition (PHC) organizers solicit proposals from any interested party for candidate password hashing schemes, to be considered for inclusion in a portfolio of schemes suitable for widespread adoption, and covering a broad range of applications.



Memory-hardness was de-facto requirement for PHC

Security

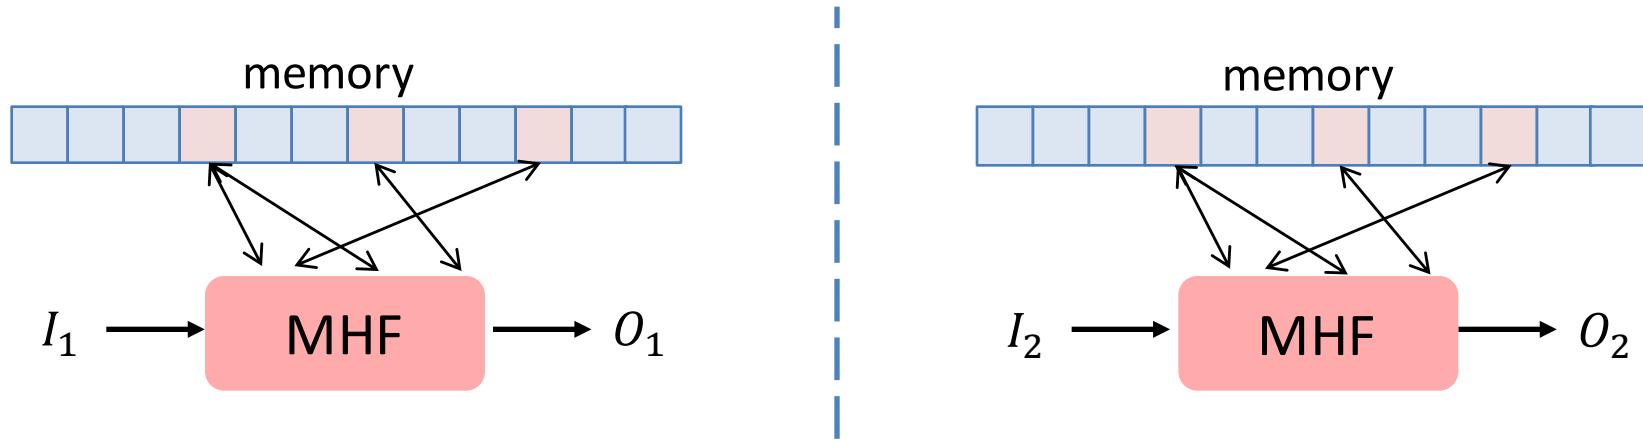
- Cryptographic security: the function should behave as a random function (random-looking output, one-way, collision resistant, immune to length extension, etc.).
- Speed-up or other efficiency improvement (e.g., in terms of memory usage per password tested) of cracking-optimized implementations (checking multiple sets of inputs in parallel, and doing so in a CPU's native code) compared to implementations intended for password validation should be minimal.
- Speed-up or other efficiency improvement (e.g., in terms of area-time product per password tested) of cracking-optimized ASIC, FPGA, and GPU implementations (checking multiple sets of inputs in parallel) compared to CPU implementations intended for password validation should be minimal.
- Resilience to side-channel attacks (timing attacks, leakages, etc.). In particular, information should not leak on a password's length.

Many memory-hard candidates: **Argon2d**, **Argon2i**, Scrypt, Lyra2, Balloon hashing, Catena, Yescrypt,

Can we build provably memory-hard functions?

Towards optimal memory hardness

Previous provably MHFs [AS15,BCS16,ABP17] are iMHFs: data-independent memory access patterns!



Two issues raised by Alwen and Blocki:

(1) No **iMHF** achieves optimal memory hardness.

(2) Practical iMHFs are even less memory hard for parallel evaluation strategies.

Can data-dependence help? ??

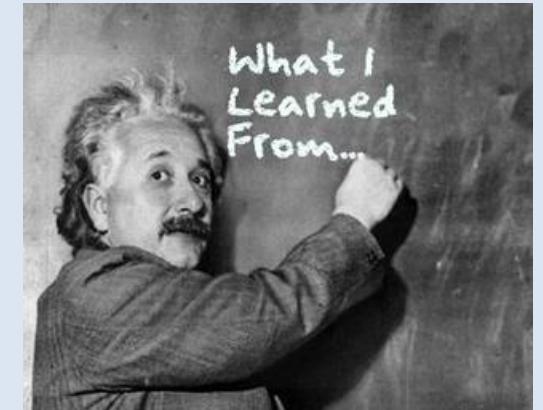


This paper: Scrypt is optimally memory hard

- Very first conjectured MHF: Proposed by **Colin Percival** in **2009**
- Used within PoWs in **Litecoin**
- Inspired the design of **Argon2d** – one of the winners of Password Hashing Competition
- Covered by RFC 7914

Take home message:

Very first example of function
with provably optimal memory
hardness.



+ it is practical, already in use, and relatively simple

Finding such proof has been a surprisingly hard problem:

- [Percival, 2009] is incorrect
- [ACKKPT16] only gave restricted result



No iMHF achieves optimal memory hardness

Roadmap

1. The Scrypt function

Definition, memory-hardness intuition, and challenges

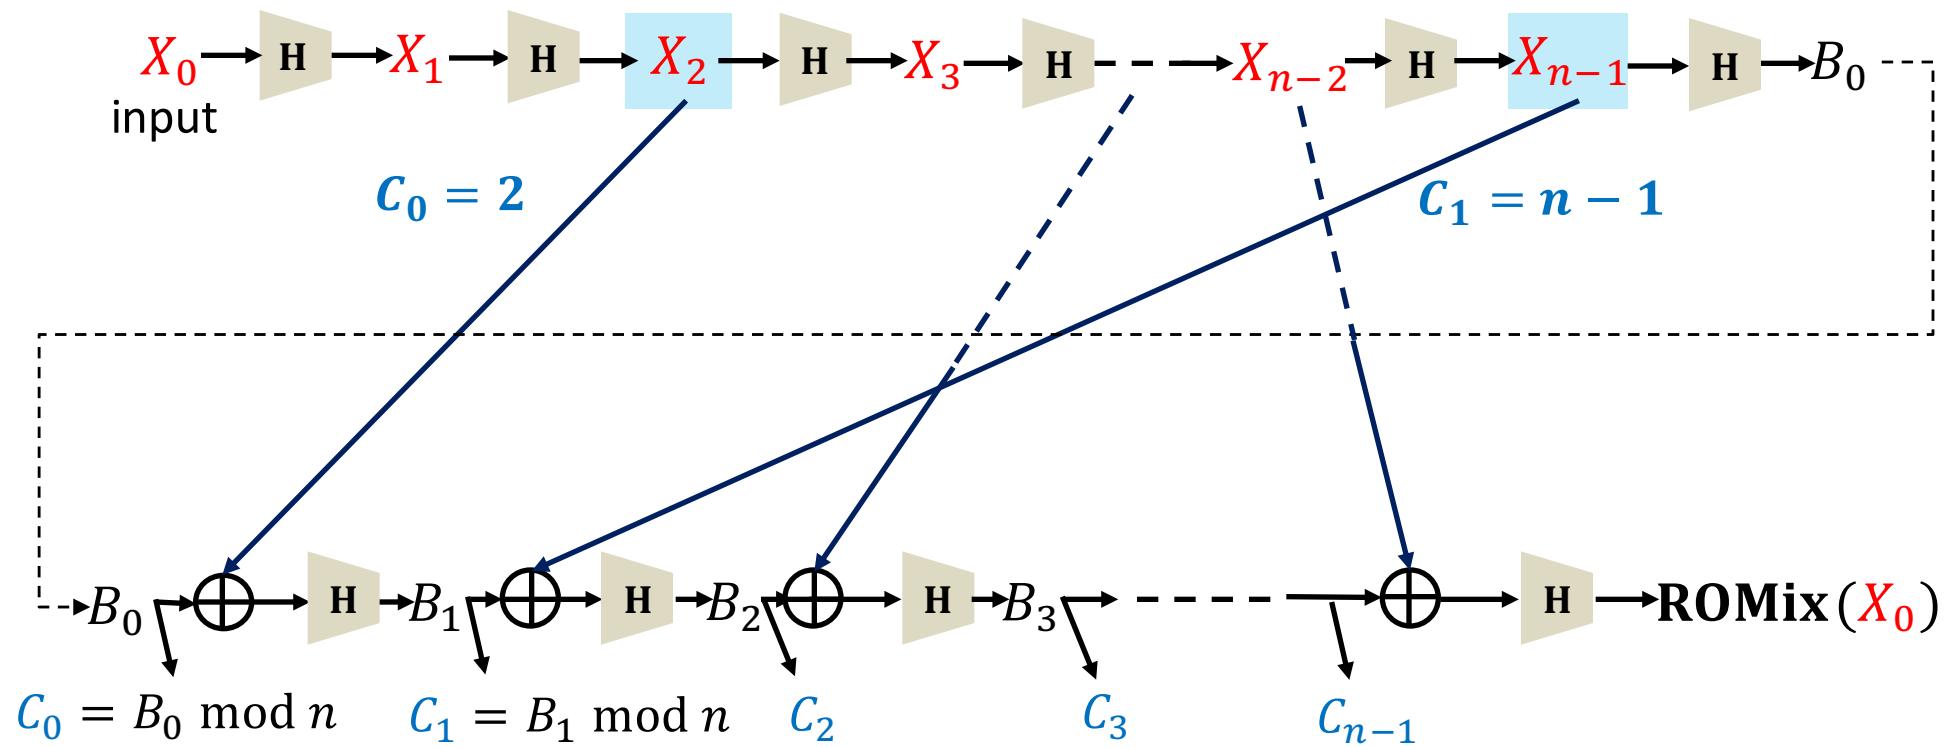
2. Optimal memory hardness of Scrypt

3. Conclusions



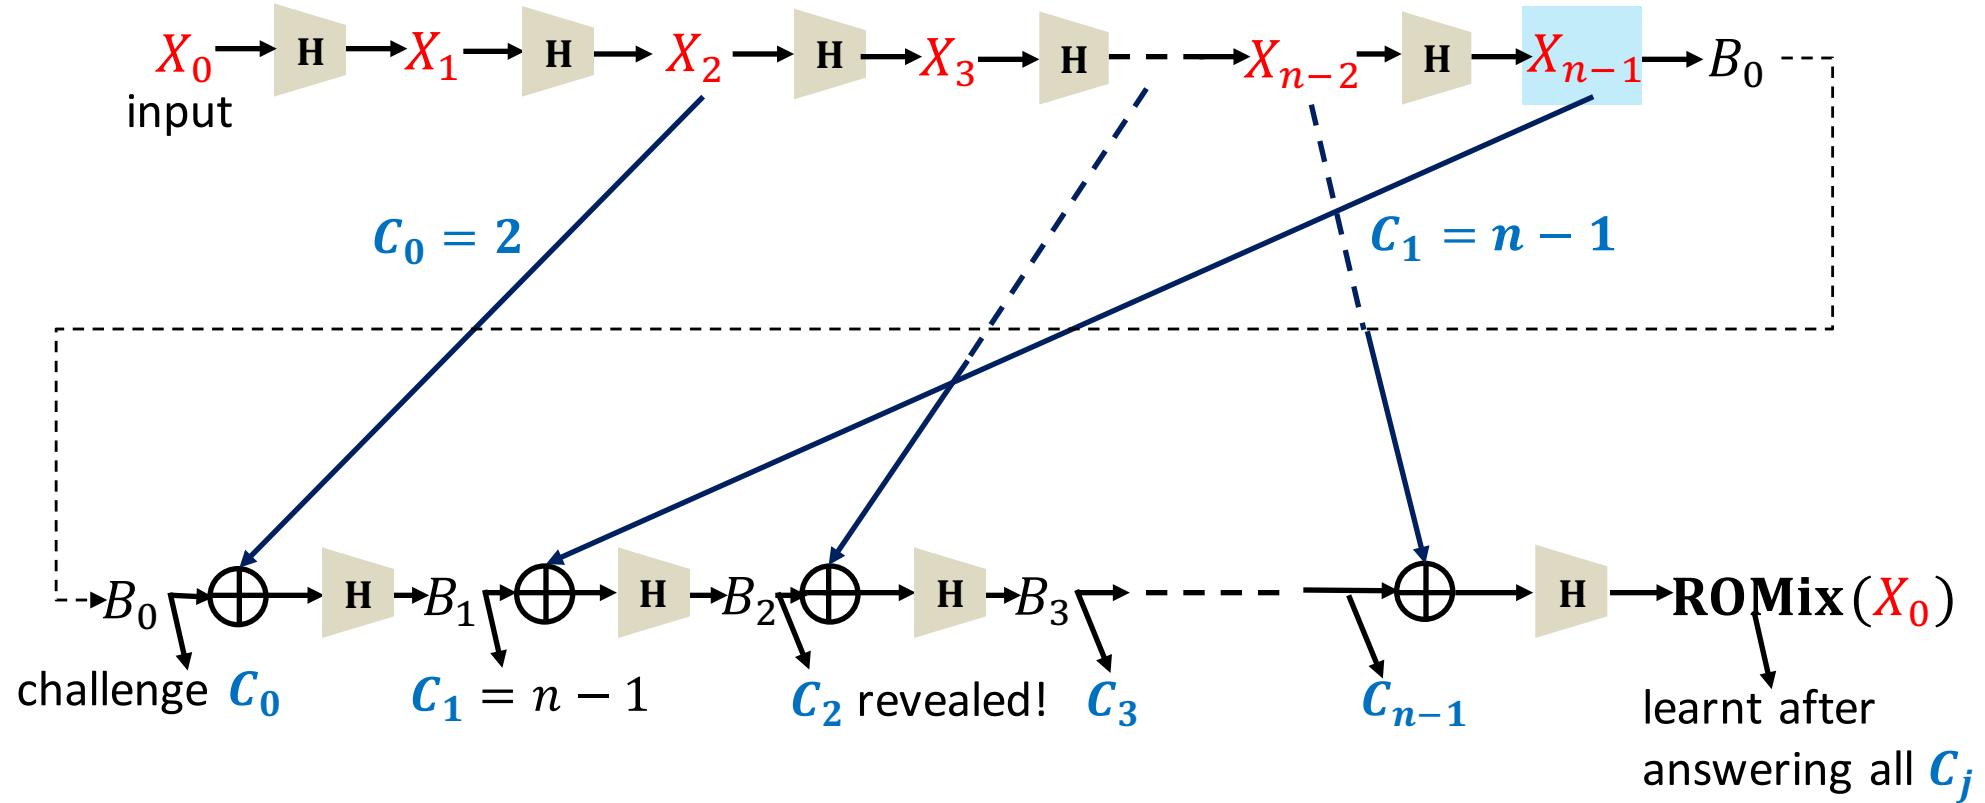
Core of Scrypt: ROMix

Modeled as a **random oracle!** \longrightarrow **H**: A Salsa20 based “hash function” with output length **w**.



n : a tunable parameter.
e.g., $n = 2^{14}$, $w = 1$ KB

ROMix: Answering challenges

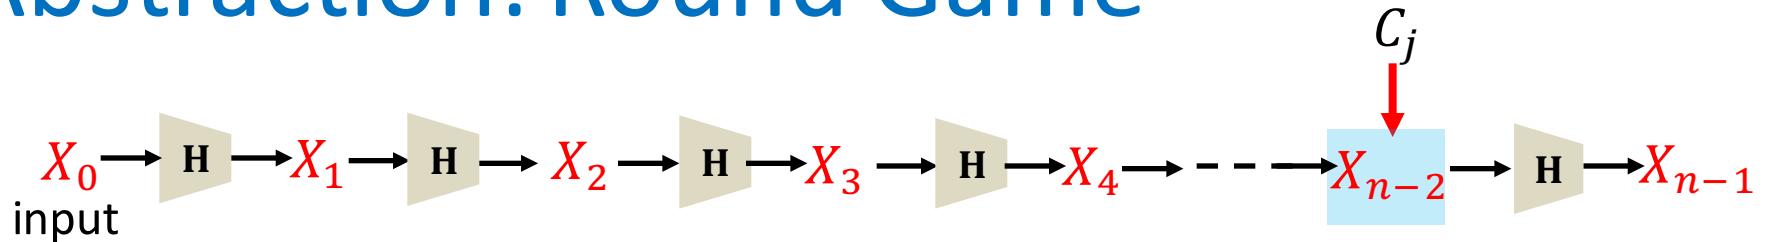


C_0, C_1, \dots, C_{n-1} **unpredictable challenges**:

1. Need to know X_{C_j} to learn C_{j+1}
2. Need to answer all challenges to complete the evaluation

Useful to abstract this!

Abstraction: Round Game



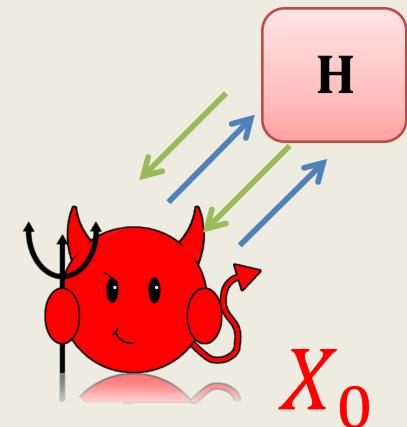
Abstract 2nd phase: challenges are ~~H~~-dependent random!

For all round $j = 0, \dots, n - 1$:



Challenger

$$C_j \leftarrow \{0, 1, \dots, n - 1\}$$



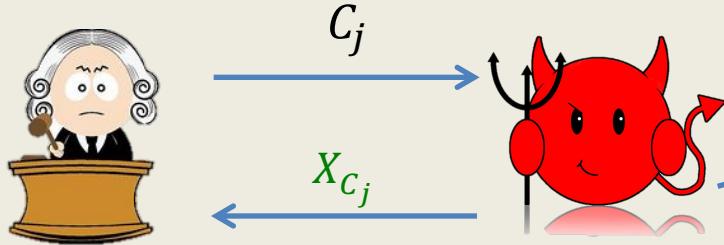
Adversary

Adversary's goal: Reduce its own **CMC** for answering all challenges!

CMC = Cumulative Memory Complexity = $\sum_{t=0}^T \text{Memory}(t)$

Round game – Naïve strategy

For round $j = 0$ to $n - 1$:



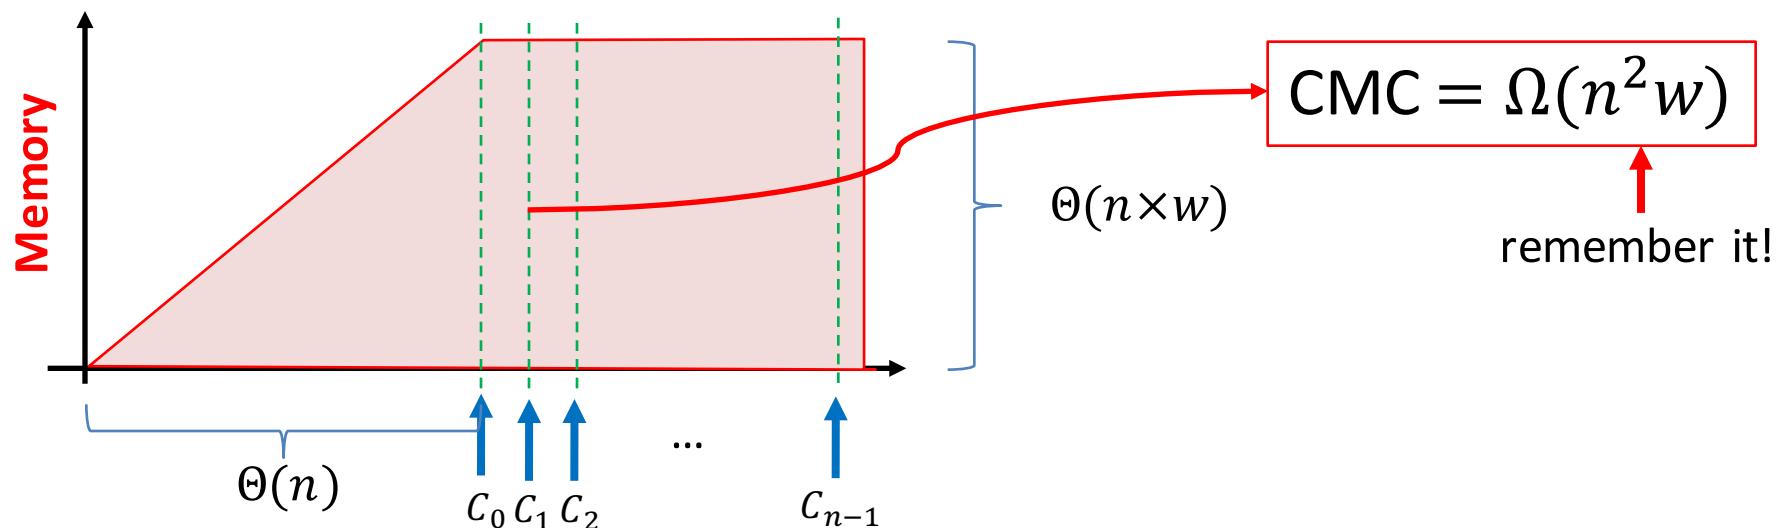
Init:

$\text{Mem}[0] \leftarrow X_0$

for $i = 1, \dots, n$ **do**

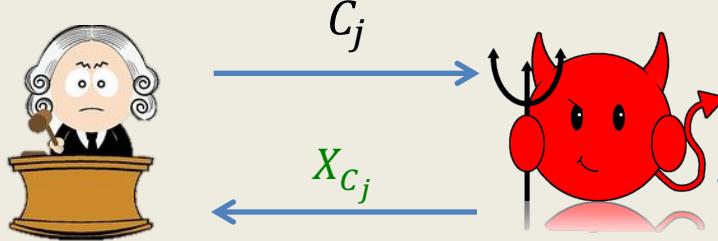
$\text{Mem}[i] \leftarrow \mathbf{H}(\text{Mem}[i - 1])$

Upon challenge C_j : **return** $\text{Mem}[C_j]$



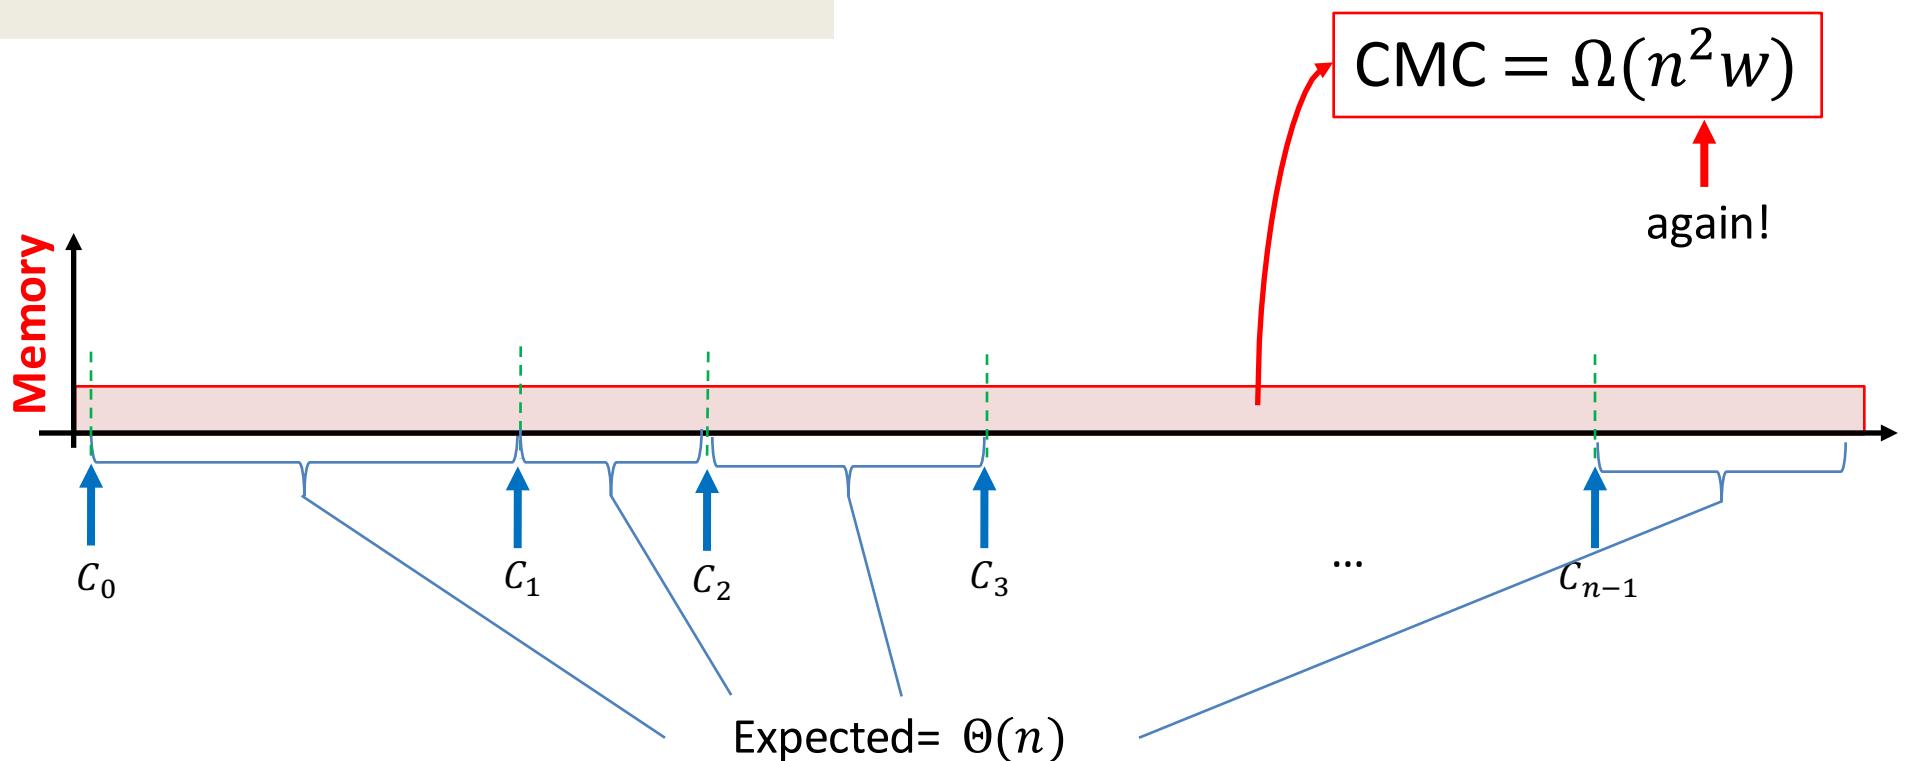
Round game – Memory-less strategy

For round $j = 0$ to $n - 1$:

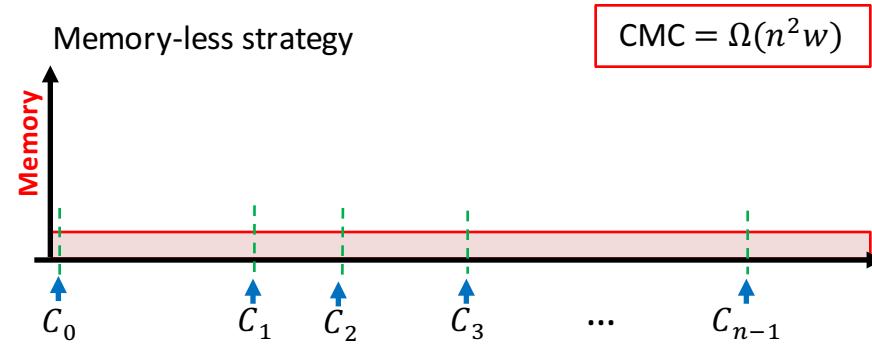
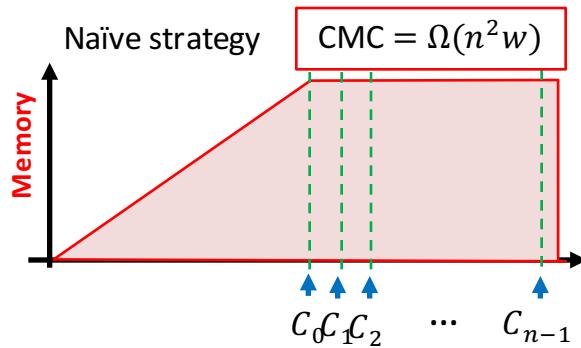


Upon challenge C_j :

```
 $X = X_0$ 
for  $i = 1, \dots, C_j$  do  $X \leftarrow H(X)$ 
return  $X$ 
```

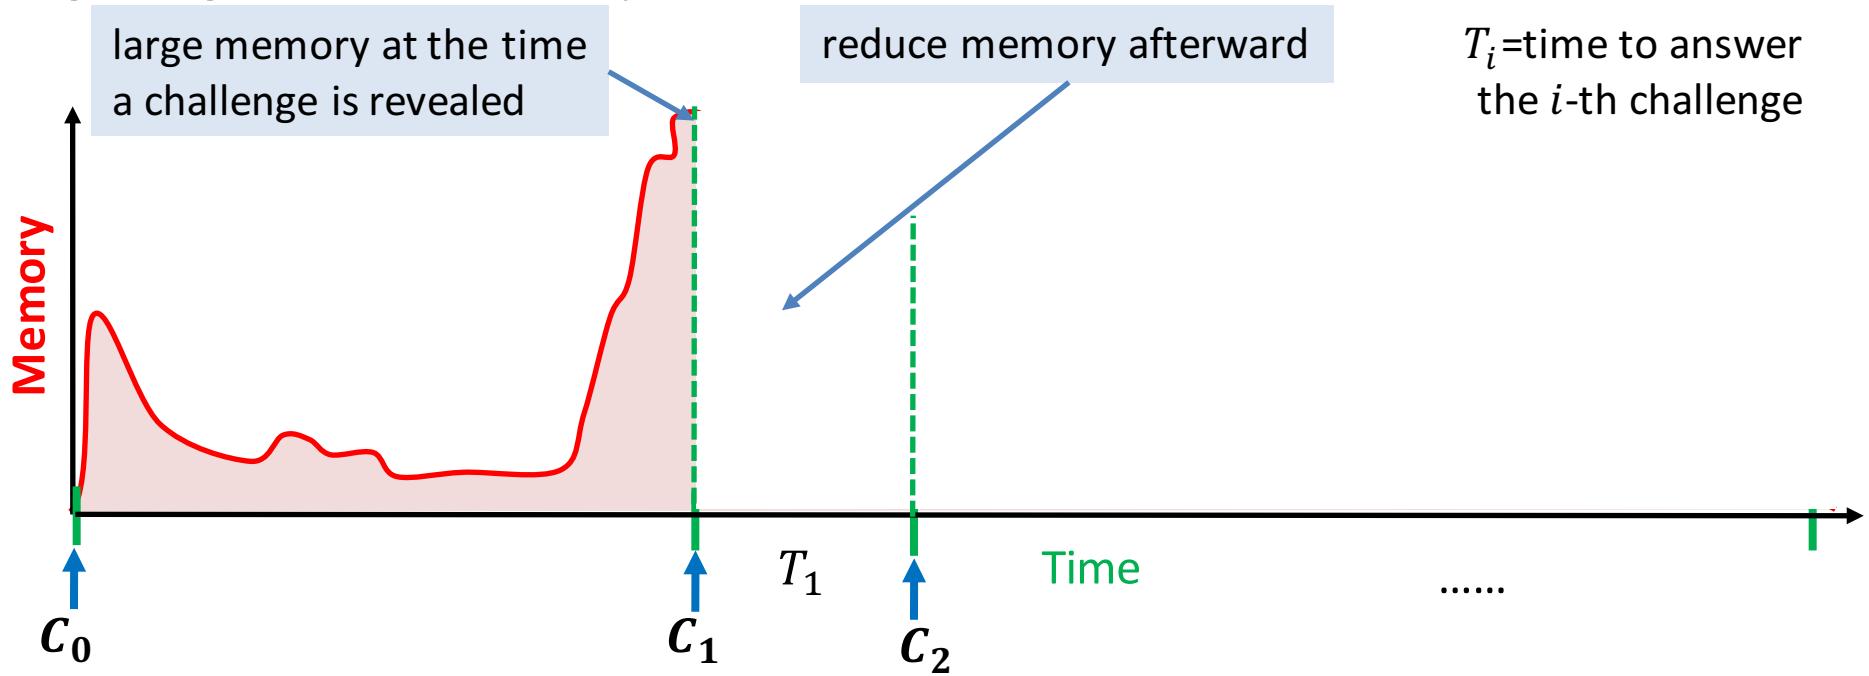


Previous two strategies are special cases: consistent memory size



More general strategy: memory consumption can vary a lot

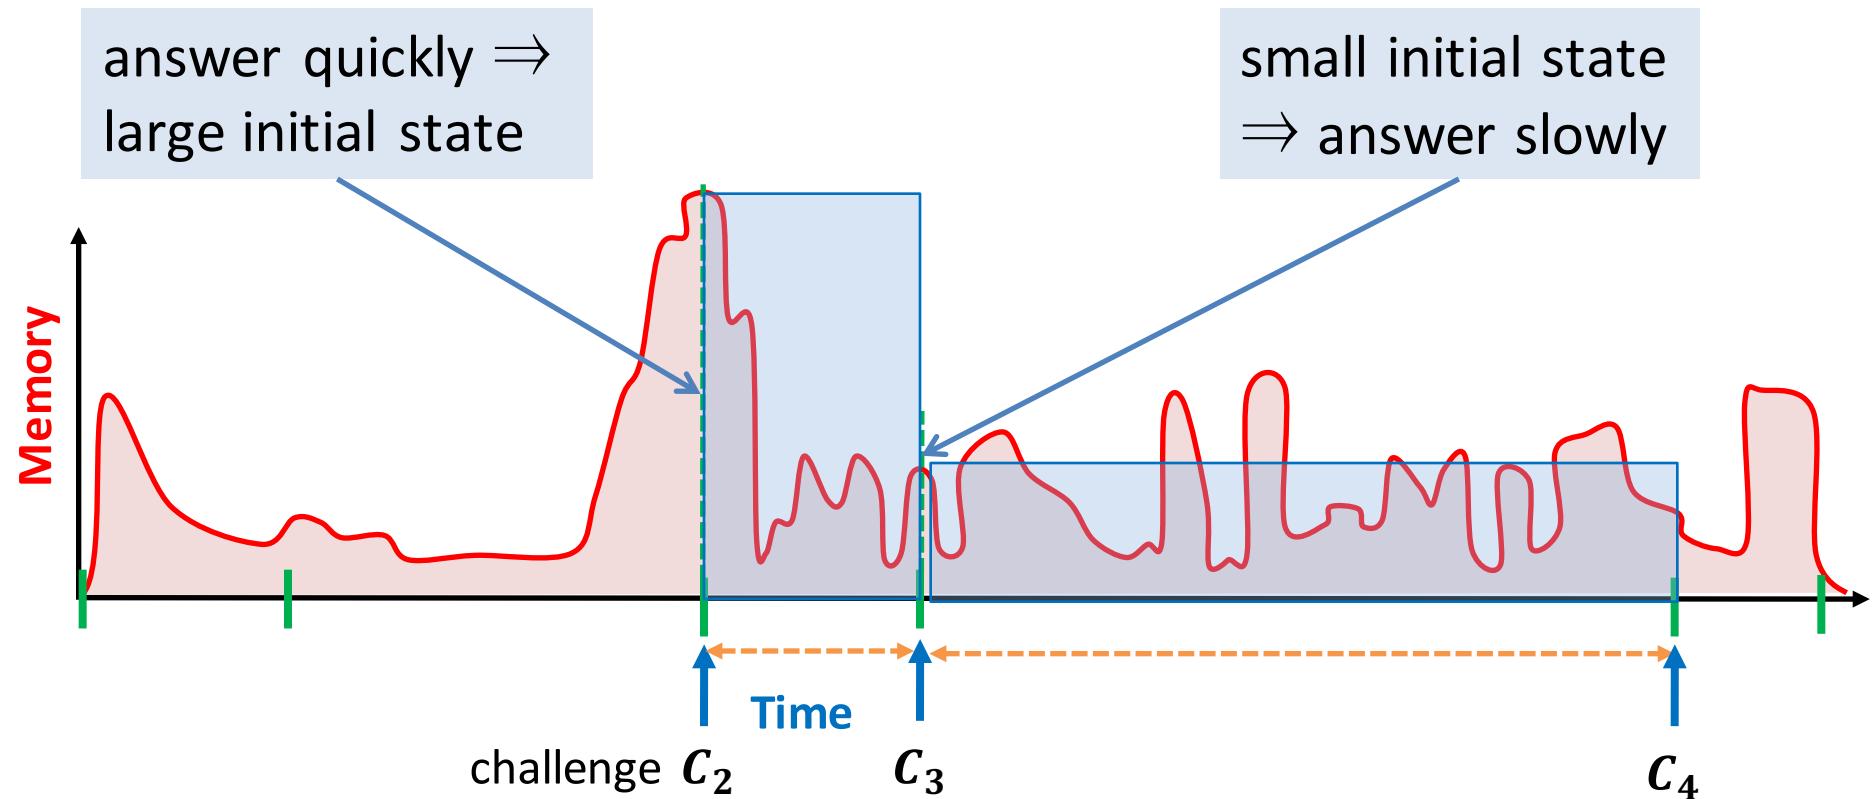
e.g., forget values, re-compute afterward



Goal: prove $CMC = \Omega(n^2 w)$!

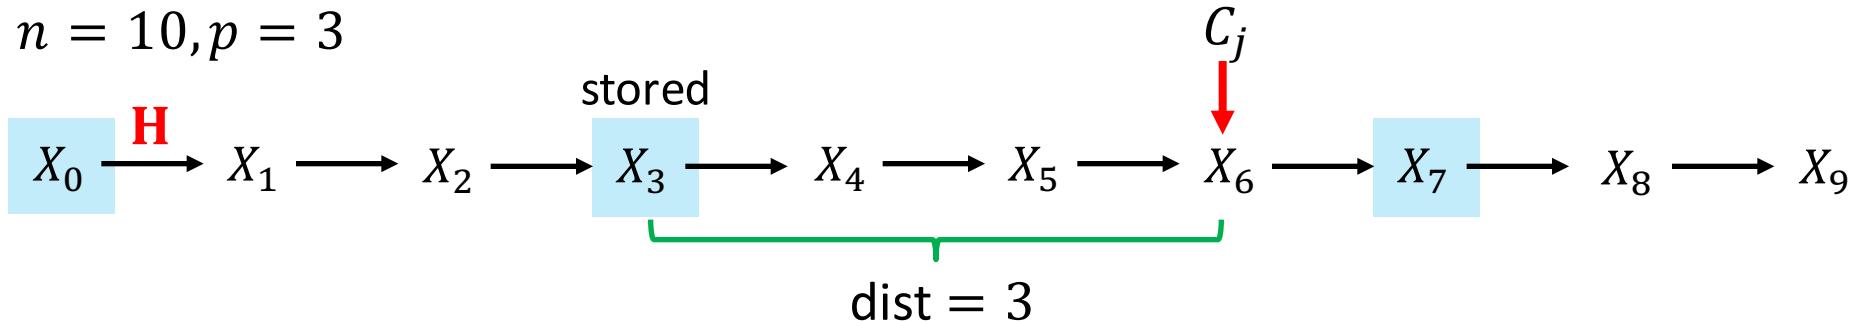
Memory hardness: intuition

Intuition: Answering challenge fast requires large state!



Single-shot memory-time trade-off

Simplifying assumption: upon learning challenge C_j , adversary only stores p of the values X_0, \dots, X_{n-1}



Fact: Avg-distance from X_{C_j} to closest stored X_i preceding X_{C_j} is $n/2p$

Regardless of parallelism, as computation of X -values is inherently sequential!

Expected time to answer the challenge is $n/2p$

$\approx |\text{memory}|$

How to translate this intuition into a memory-hardness proof for ROMix?

Three technical barriers:

1. Adversary stores arbitrary information

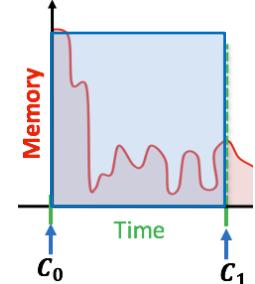
e.g., XOR of X_i values, halves of X_i , reconstruct information adaptively on challenges, etc.

[ACKKPT16] considered restricted strategies and exhibited round games where general storing strategies can help!

2. Memory variation during computation
single-shot memory-time trade-off not enough!

[ACKKPT16] only shows CMC = $\Omega\left(\frac{n^2 w}{\log^2(n)}\right)$

**Focus on
1 and 2**



3. H-dependent challenges, as opposed to truly random **see the paper!**

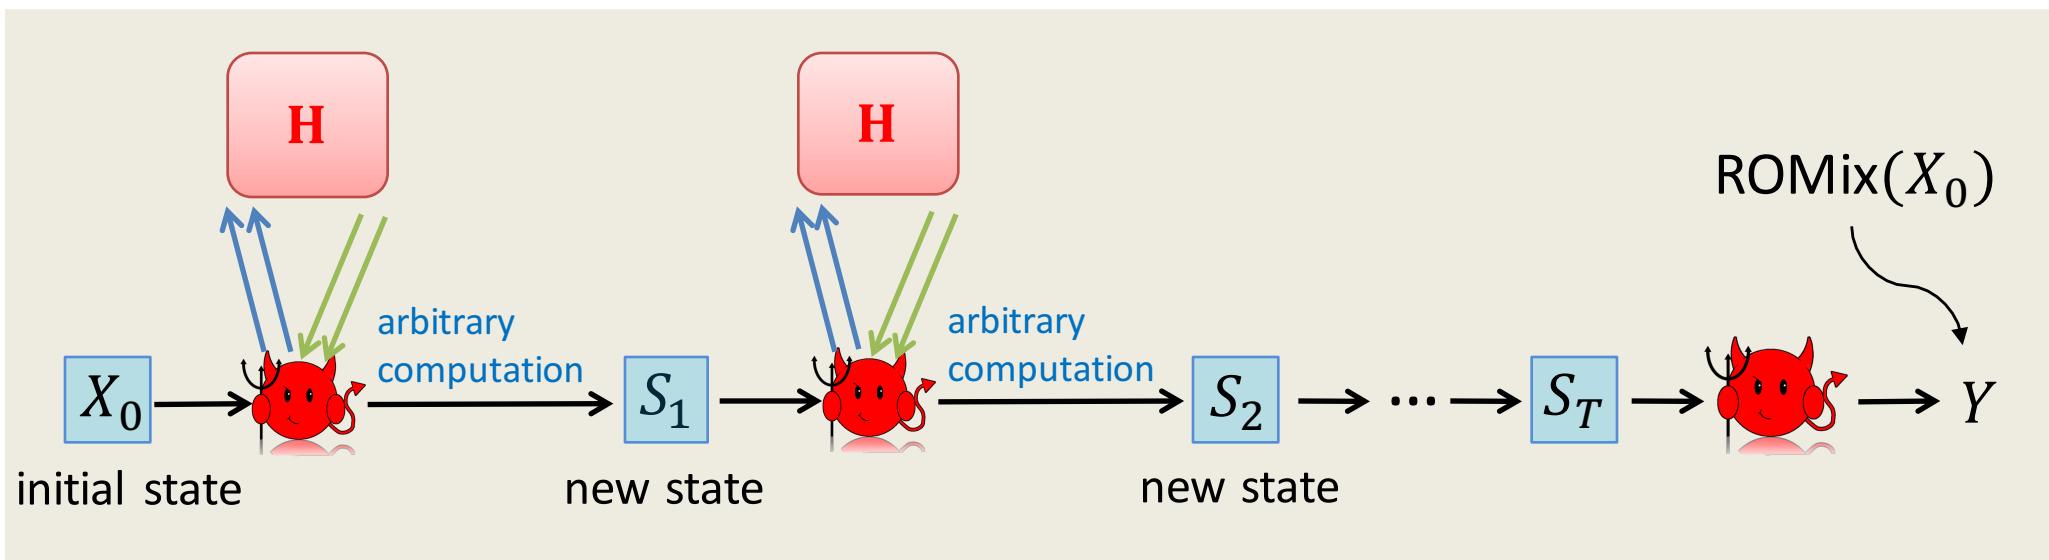
Roadmap

1. The Scrypt function
2. Optimal memory hardness of Scrypt
Model, theorem, and proof approach
3. Conclusions



The parallel random oracle model

[Alwen and Serbinenko, STOC '15]



At each step: Adv asks one batch of parallel H queries
+ performs unbounded computation

Goal of adv: minimize $\text{CMC} = \sum_{i=1}^T |S_i|$

Main Theorem.

For any adversary **A** evaluating **ROMix**,

$$\text{CMC}(\mathbf{A}) \geq \frac{1}{25} \cdot n^2 \cdot (w - 4 \cdot \log(n))$$

w/ overwhelming probability over the choice of **H**.

The $4\log(n)$ loss is inherent in the proof.

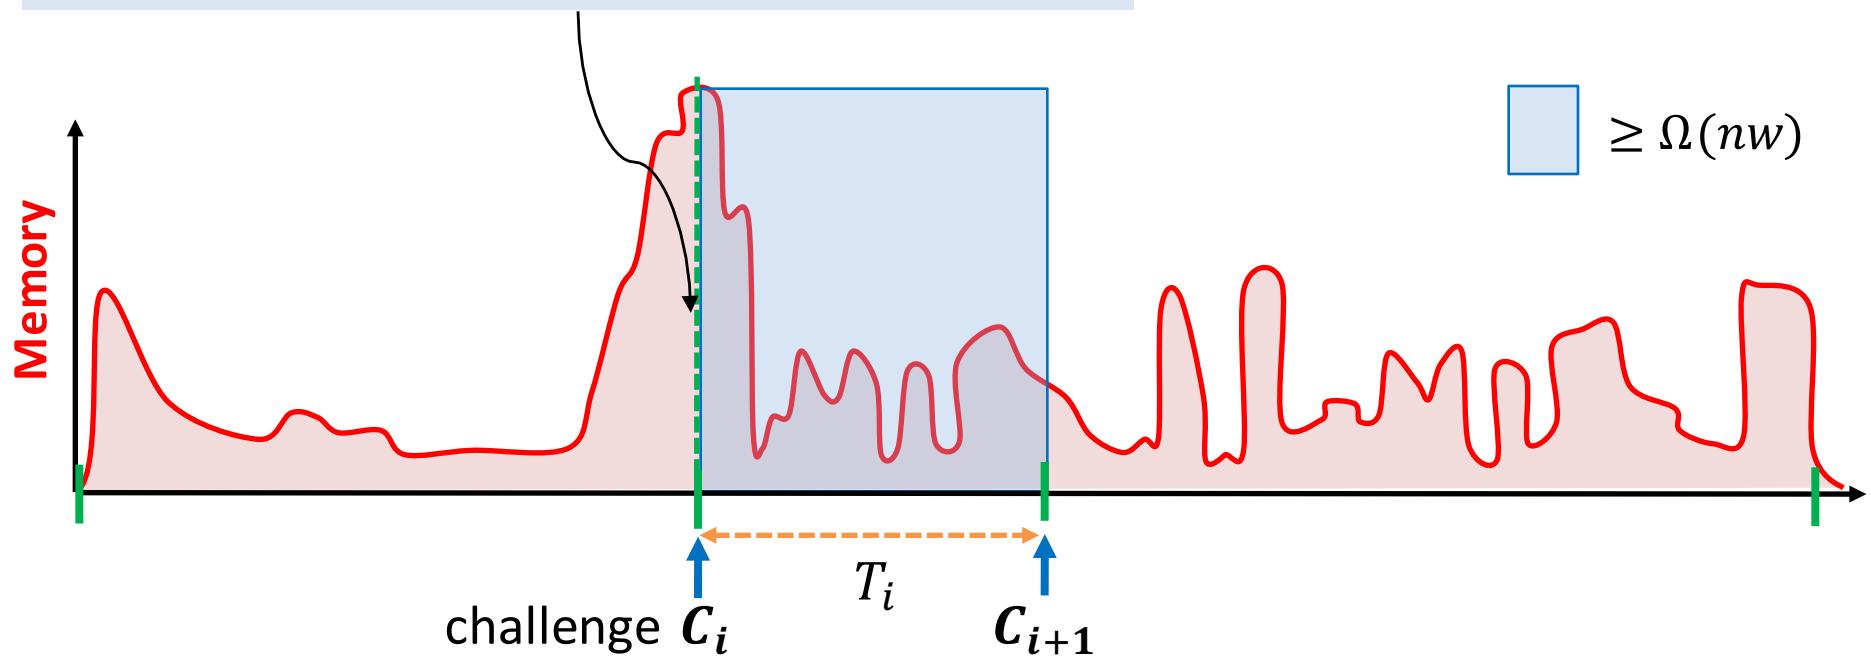
$\Omega(n^2 w)$ clearly best possible for any construction making n queries to **H**.

Naïve strategy: Make n calls, remember all outputs

Proof strategy: step 1

Green (memory usage at this step)
is inversely proportional to orange (T_i)

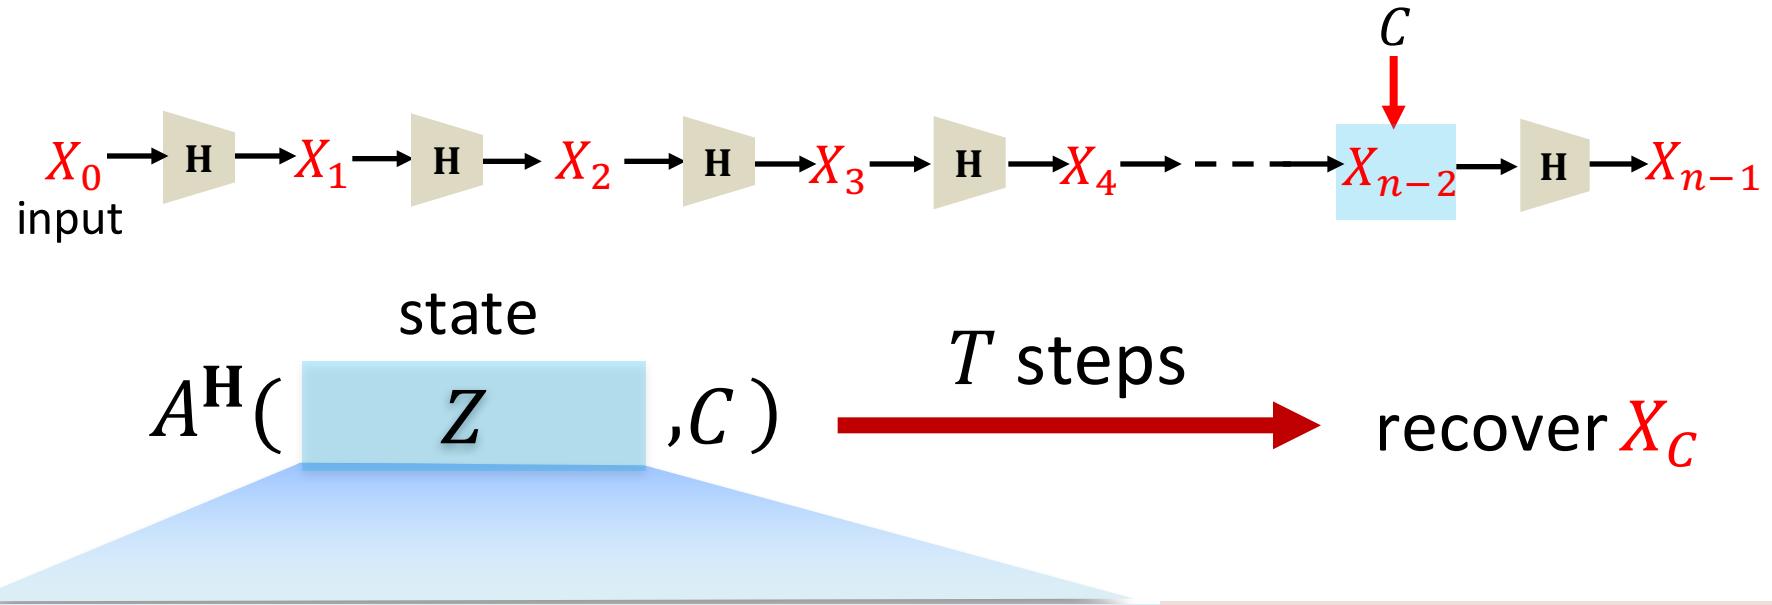
T_i =time to answer the
 i -th challenge



Memory-time trade-off \Rightarrow lower bound on memory

The memory-time trade-off holds true for adv ~~storing X-values~~
even if the adv stores arbitrary information!

Single-shot memory-time trade-off

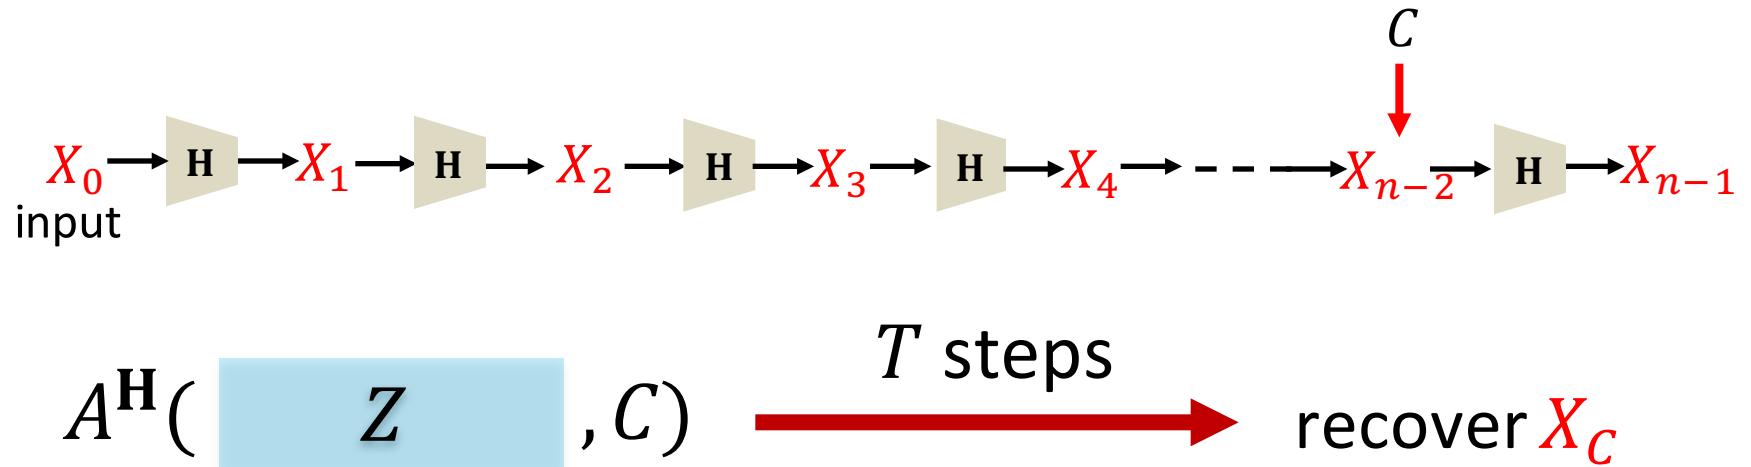


Z : arbitrary computation on H -outputs Goal: Lower bound $|Z|$
• E.g., pre-computation of H 's entries, as function of T and n
XOR of $\{X_i\}$ values, halves of X_i

[ACKKPT16]: computation on H -outputs can help in some round games

[This result]: computation on H -outputs cannot help for Scrypt!

Single-shot memory-time trade-off



Lemma. For all A , for most \mathbf{H} , if $|Z| \approx pw$ bits

$$\Pr_C \left[T > \frac{n}{2p} \right] > \frac{1}{2}$$

Lemma. For all A , for most \mathbf{H} , if $|Z| \approx pw$ bits

$$\Pr_C \left[T > \frac{n}{2p} \right] > \frac{1}{2}$$

Proof idea:

If adversary $A^{\mathbf{H}}(Z, C)$ answers too fast for most challenges C



$A^{\mathbf{H}}(Z, C)$ can output or query many X_i values w/o querying \mathbf{H} first

Can compress the oracle \mathbf{H} using state Z

Cannot be true for too many \mathbf{H} : random oracle is incompressible

[Dwork, Naor and Wee, Crypto'05], [Alwen and Serbinenko, STOC '15]

Proof strategy: step 2

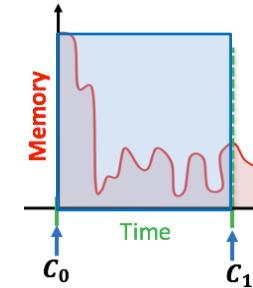
Technical barriers:

1. Adversary stores arbitrary information



Single-shot memory-time trade-off for arbitrary adv

2. Memory variation during computation



Single-shot
memory-time
trade-off



Optimal CMC
lower bound for
the round game

Generalize

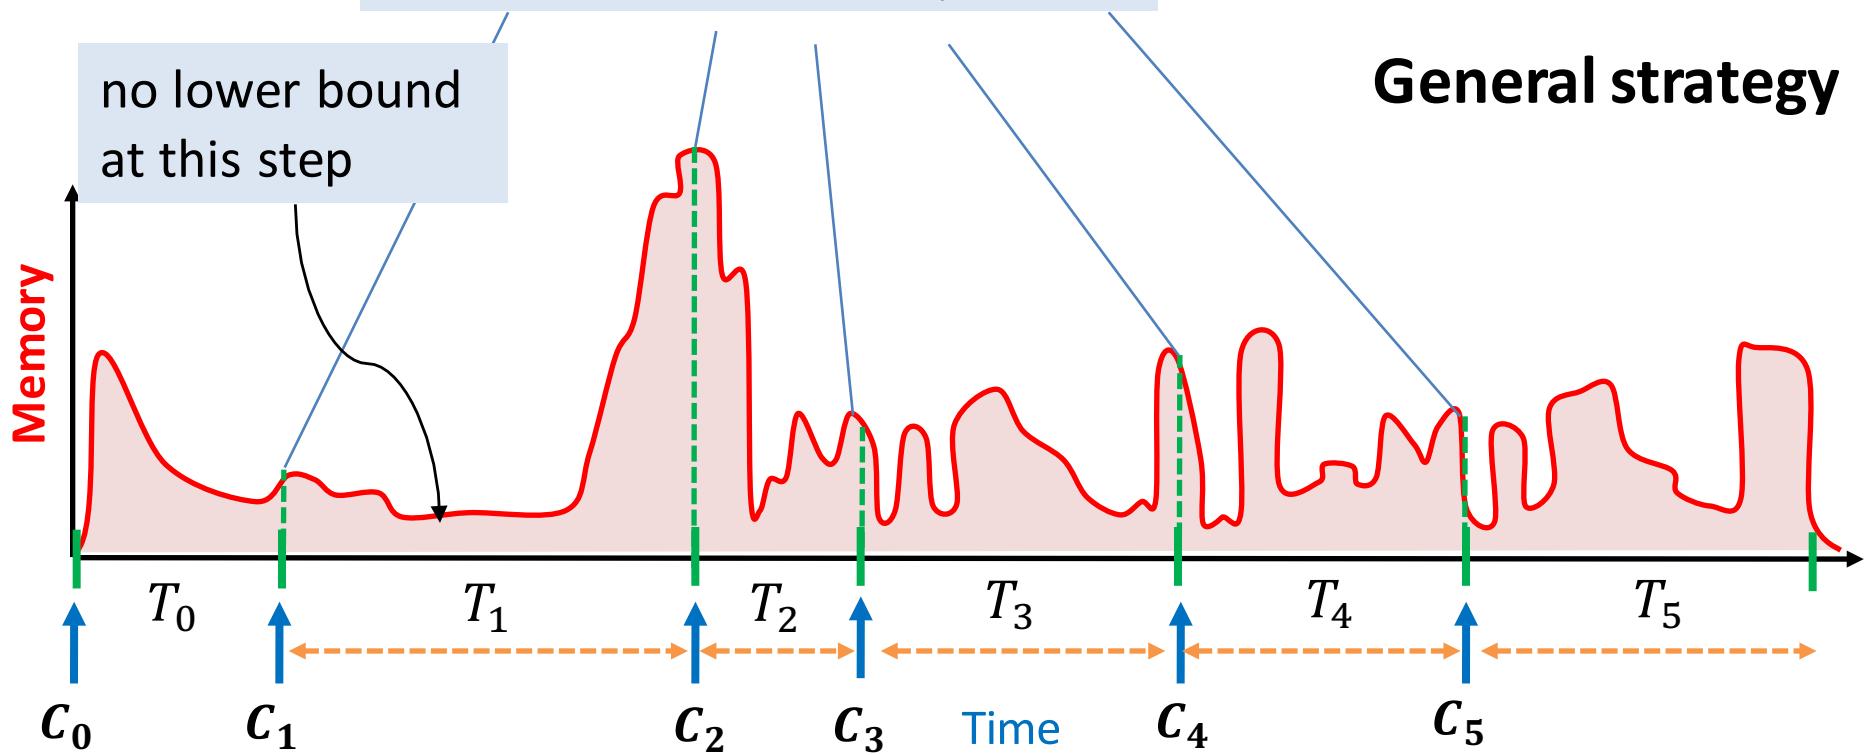
CMC lower bound

Lemma. $\Pr_C \left[p > \frac{n}{2T} \right] > \frac{1}{2}$

when learning the i -th challenge

$$|\text{memory}| \geq \frac{nw}{2T_i}$$

T_i =time to answer the i -th challenge



memory-time trade-off \Rightarrow memory lower bound
for the step right before the challenge is revealed

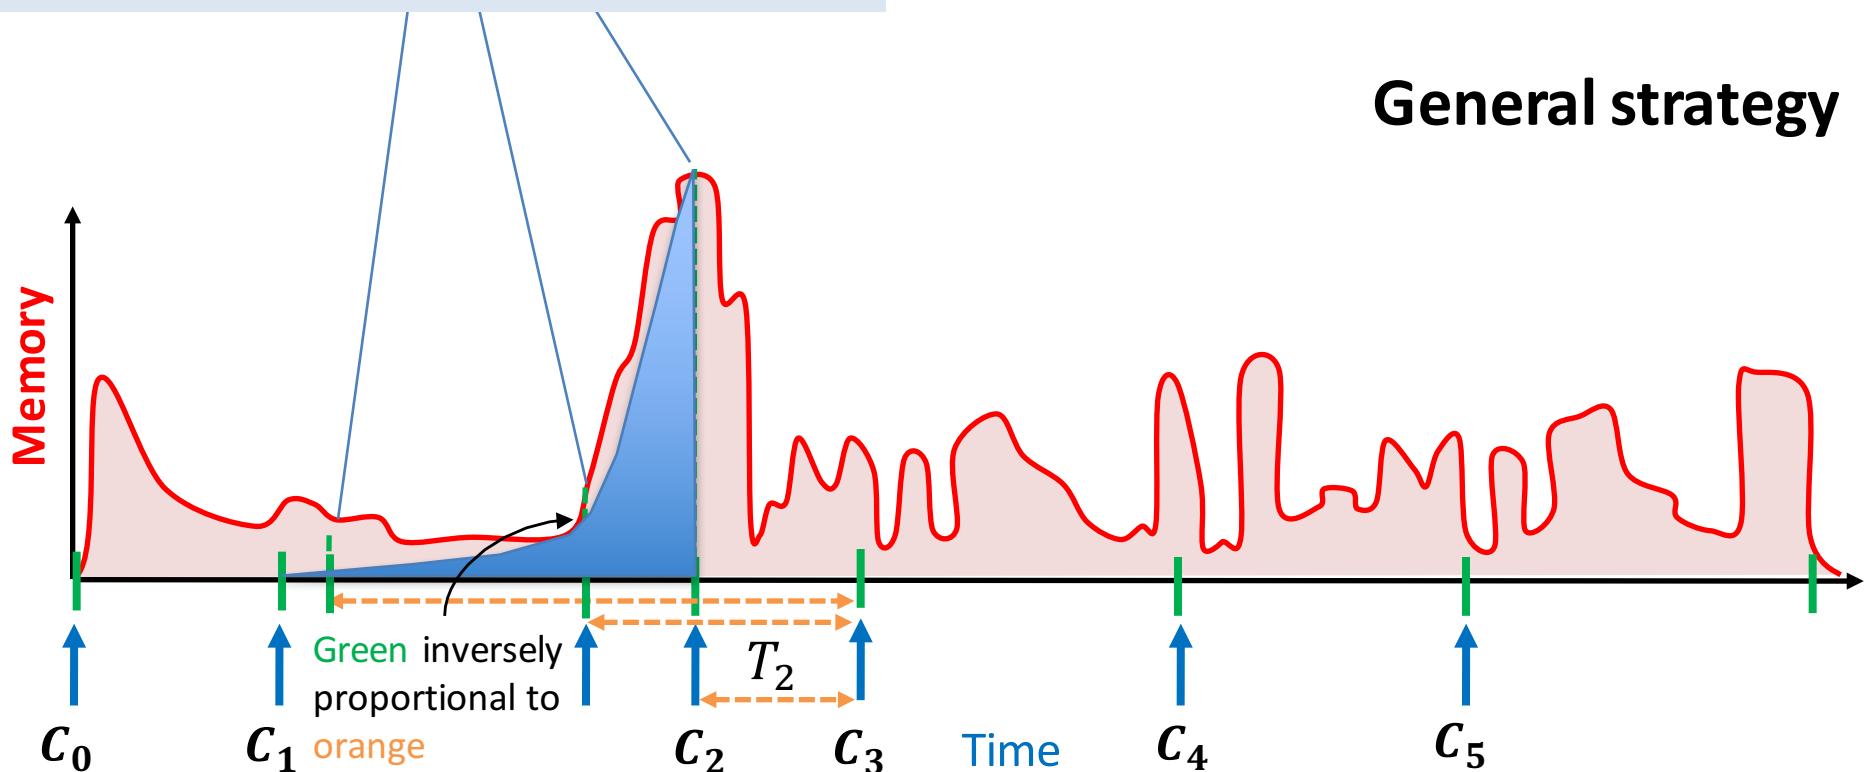
CMC = ???

CMC lower bound

Lemma. $\Pr_C \left[p > \frac{n}{2T} \right] > \frac{1}{2}$

mem at every step \geq funcs of n and
time to answer the next challenge

T_i =time to answer the i -th challenge



Similar trade-off holds for every step before challenge is revealed

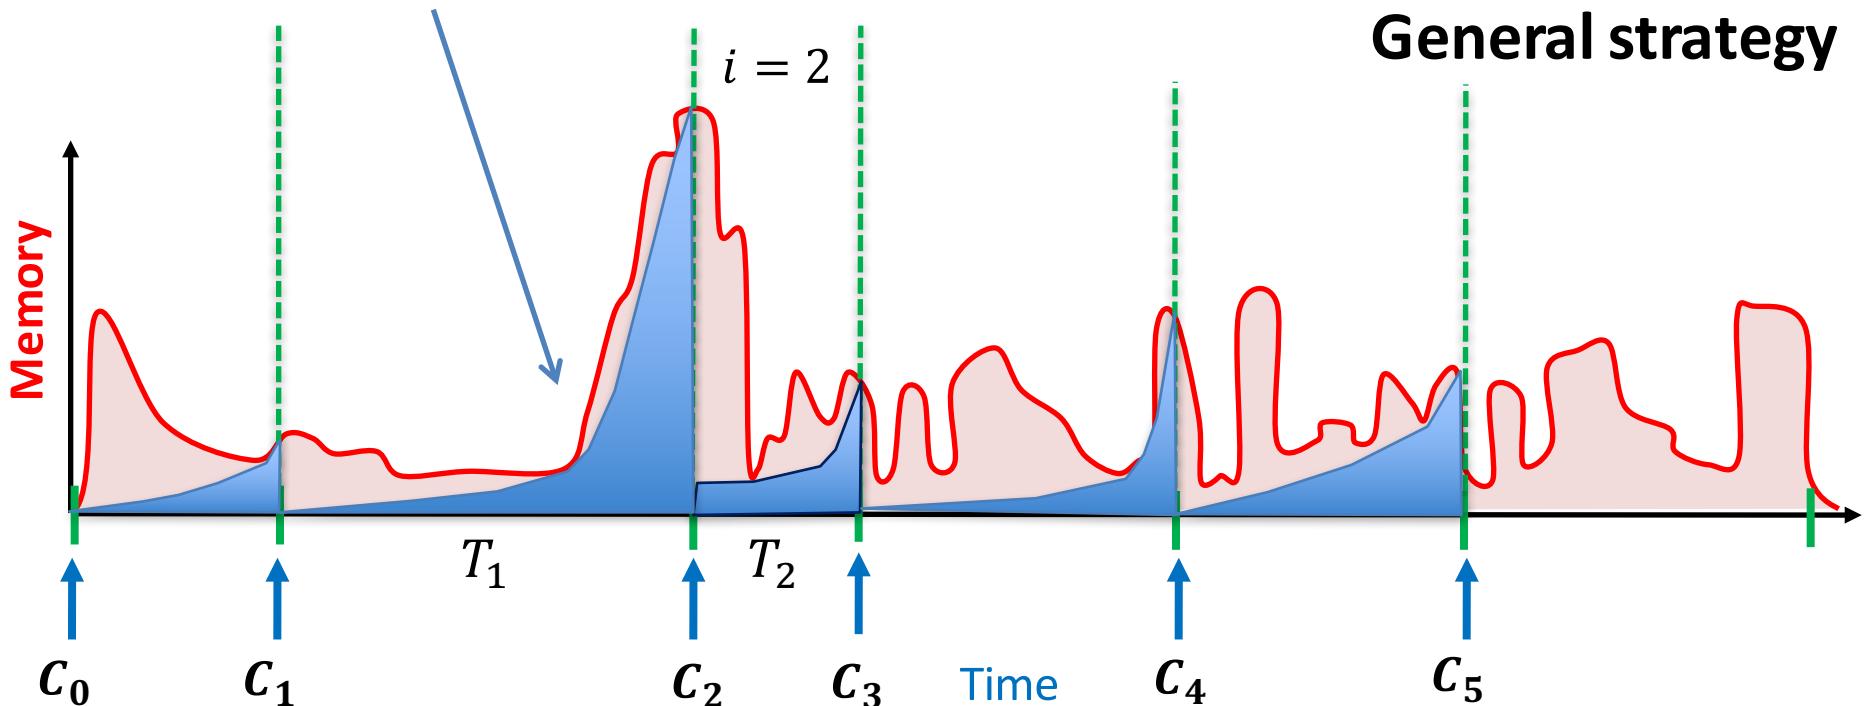
CMC lower bound

Lemma. $\Pr_C \left[p > \frac{n}{2T} \right] > \frac{1}{2}$

During round $i - 1$:

$$\text{Sum of memory} \geq \frac{nw}{2} \ln \left(1 + \frac{T_{i-1}}{T_i} \right)$$

T_i = time to answer the i -th challenge



By adding lower bounds over rounds from 0 to $n - 1$, we have $\text{CMC} = \Omega(n^2 w)$

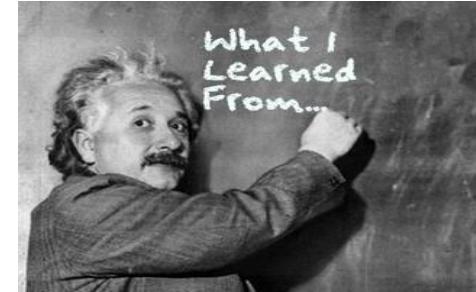
$$= \Omega(n^2 w)$$

Roadmap

1. The Scrypt function
2. Optimal memory hardness of Scrypt
3. Conclusions



Summary



- Script is **maximally** memory hard
 - First optimal memory-hardness proof.
 - Validates a practical MHF design.
- Open problem
 - Optimal memory hardness proof for Argon2d?

Thank you! – Merci!

<https://eprint.iacr.org/2016/989>